

## HW 4

[1]. I'd transform degree to create our treatment variable d. What would you do and why?

- Transformation:  $\log(\text{degree}+1)$
- In this way, I can normalize distribution, resolving the skewness in the degree data. The distribution of d becomes closer to the normal gaussian distribution and suitable for fitting linear regression. Instead of simple log transformation,  $\log(x + 1)$  is used for values that contain 0. This enable us to avoid log transformation returning negative infinity.

```
hh <- read.csv("microfi_households.csv", row.names="hh")
hh$village <- factor(hh$village)

## We'll kick off with a bunch of network stuff.
edges <- read.table("microfi_edges.txt", colClasses="character")
## edges holds connections between the household ids
hhnet <- graph.edgelist(as.matrix(edges))
hhnet <- as.undirected(hhnet) # two-way connections.

## igraph is all about plotting.
V(hhnet) ## our 8000+ household vertices

## + 8182/8182 vertices, named, from c06095e:
##   [1] 1002 1001 1020 1042 1053 1163 1003 1004 1026 1029 1076
##   [12] 1159 1106 1031 1048 1081 1006 1005 1008 1016 1021 1024
##   [23] 1089 1103 1007 1019 1155 1015 1040 1044 1045 1078 1088
##   [34] 1110 1115 1140 1145 1009 1018 1060 1064 1073 1153 1067
##   [45] 1099 1010 1162 1012 1143 1013 1023 1028 1034 1065 1117
##   [56] 1139 1154 1157 1173 1014 1068 1071 1148 1017 1036 1062
##   [67] 1112 1118 1120 1129 1134 1165 1183 1126 1122 1049 1058
##   [78] 1093 1108 1114 1119 1022 1043 1079 1033 1102 1104 1105
##   [89] 1152 1169 1171 1025 1027 1147 1032 1035 1037 1039 1041
##  [100] 1113 1174 1069 1116 1132 1178 1146 1080 1086 1101 1172
## + ... omitted several vertices


## Each vertex (node) has some attributes, and we can add more.
V(hhnet)$village <- as.character(hh[V(hhnet), 'village'])
## we'll color them by village membership
vilcol <- rainbow(nlevels(hh$village))
names(vilcol) <- levels(hh$village)
V(hhnet)$color = vilcol[V(hhnet)$village]
## drop HH labels from plot
V(hhnet)$label=NA

# graph plots try to force distances proportional to connectivity
# imagine nodes connected by elastic bands that you are pulling apart
# The graphs can take a very long time, but I've found
```


```
# edge.curved=FALSE speeds things up a lot. Not sure why.

## we'll use induced.subgraph and plot a couple villages
village1 <- induced.subgraph(hhnet, v=which(V(hhnet)$village=="1"))
village33 <- induced.subgraph(hhnet, v=which(V(hhnet)$village=="33"))

# vertex.size=3 is small. default is 15
plot(village1, vertex.size=3, edge.curved=FALSE)
```



```
plot(village33, vertex.size=3, edge.curved=FALSE)
```




```
##### now, on to your homework stuff
## match id's; I call these 'zebras' because they are like crosswalks
zebra <- match(rownames(hh), V(hhnet)$name)

## calculate the `degree' of each hh:
## number of commerce/friend/family connections
degree <- degree(hhnet)[zebra]
names(degree) <- rownames(hh)
degree[is.na(degree)] <- 0 # unconnected houses, not in our graph


## if you run a full glm, it takes forever and is an overfit mess
```

```
# > summary(full <- glm(loan ~ degree + .^2, data=hh, family="binomial"))  
# Warning messages:  
# 1: glm.fit: algorithm did not converge  
# 2: glm.fit: fitted probabilities numerically 0 or 1 occurred
```

```
hist(degree)
```



```
hist(log(degree+1))
```



```
d<-log(degree+1)
```

[2]. Build a model to predict  $d$  from  $x$ , our controls. Comment on how tight the fit is, and what that implies for estimation of a treatment effect.

- $R^2$  using normal regression: 0.08223472
- $R^2$  using LASSO: 0.08187873
- Since  $R^2$  is quite low, this indicates that the part of  $d$  that can be predicted with  $x$ 's isn't very high. This implies that the treatment effect, when estimated without the controls will not be overestimated by too much, given that the controls we have are enough to account for the confounding effect.

```
#without LASSO
reg_dx <- glm(d ~ .-loan, data=hh)
summary(reg_dx)

1-5826.8/6348.9

## [1] 0.08223472

#With LASSO
x = sparse.model.matrix(~.-loan, data=hh)[, -1]
treat <- gamlr(x,d,lambda.min.ratio=1e-4)
dhat <- predict(treat, x, type="response")
cor(drop(dhat),d)^2

## [1] 0.08187873
```

[3]. Use predictions from [2] in an estimator for effect of  $d$  on loan.

```
dhat<-reg_dx$fitted.values

causal <- gamlr(cBind(d,dhat,x),hh$loan,free=2,lmr=1e-4)

coef(causal)["d",]

## [1] 0.01803462
```

[4]. Compare the results from [3] to those from a straight (naïve) lasso for loan on  $d$  and  $x$ . Explain why they are similar or different.

- The results are similar because the dependent part of  $d$  wasn't large, so including  $x$  in the naïve LASSO would not change the estimate for the coefficient of  $d$  by too much.

```
naive <- gamlr(cBind(d,x),hh$loan)

coef(naive)["d",]

## [1] 0.01868003
```

[5]. Bootstrap your estimator from [3] and describe the uncertainty.

- The standard error for the estimator is 0.00376866. This represents the uncertainty of the estimator because given  $\pm$  one standard deviation from the estimator, about 68.3% of the time, the true value of the measured quantity falls within the stated uncertainty range.

```
y <- hh$loan
n <- nrow(x)

gamb <- c() # empty gamma

for(b in 1:20){
  ## create a matrix of resampled indices

  ib <- sample(1:n, n, replace=TRUE)

  ## create the resampled data

  xb <- x[ib,]
  db <- d[ib]
  yb <- y[ib]

  ## run the treatment regression

  treatb <- gamlr(xb,db,lambda.min.ratio=1e-3)
  dhatb <- predict(treatb, xb, type="response")
  fitb <- gamlr(cBind(db,dhatb,xb),yb,free=2)

  gamb <- c(gamb,coef(fitb)["db",])

  print(b)
}

## [1] 1
## [1] 2
## [1] 3
## [1] 4
## [1] 5
## [1] 6
## [1] 7
## [1] 8
## [1] 9
## [1] 10
## [1] 11
## [1] 12
## [1] 13
## [1] 14
## [1] 15
## [1] 16
## [1] 17
## [1] 18
```

```
## [1] 19
## [1] 20

summary(gamb)

##      Min.   1st Qu.   Median     Mean   3rd Qu.    Max.
## 0.009246 0.014755 0.018455 0.017108 0.020190 0.021830

coef(causal)["d",]/sd(gamb)

## [1] 4.78542

se<-sd(gamb)
se

## [1] 0.00376866

{hist(gamb)
  abline(v=quantile(gamb,0.025),col=3,lwd=2)
  abline(v=quantile(gamb,0.975),col=3,lwd=2)}
```

